

Bridging Vocabulary Mismatch in Industrial RAG Systems

Adaptive Multi-Relational Knowledge Routing, Path-Aware Neural Reranking, and Deterministic Safety Guardrails in Smart Manufacturing

Topic 1 Deliverable & Doctoral Candidacy Proposal

Academic Chair:

Management Intelligence Services (MIS)

Institute of Information Systems, FAU Erlangen-Nürnberg

Supervisors:

- **Prof. Dr. Freimut Bodendorf** (Head of Chair)
- **Yannick Rank, M.Sc.** (Advisor, Knowledge Management)

Project Framework:

In cooperation with **REHAU Industries**

Focus Areas:

- ▶ Shopfloor Terminology Alignment
- ▶ Negative Transfer / Fusion Dilution
- ▶ Neuro-Symbolic Graph Provenance
- ▶ Zero-Tolerance Safety Verification

Presentation Duration: 20 Minutes Presentation · 10 Minutes Academic Defense & PhD Discussion

Motivation: The Industrial Shopfloor Reality

Bridging Operator Phrasing and Standardized Engineering Knowledge

The Frontline Operator Language

Operators query systems under time pressure using **colloquial phrasing**:

- ▶ *"How do I prevent edge chipping on wood-look panels?"*
- ▶ *"Can I clean shade panels with standard glass cleaner?"*
- ▶ *"The back of the laminate is not holding adhesive."*

→ High domain nuance, zero formal catalog keywords.

Authoritative Engineering Manuals

Technical documentation, DIN sheets, and manuals use **standard terms**:

- ▶ *"Sacrificial underlayment below RAUVISIO ingrain..."*
- ▶ *"Alcohol/ammonia strictly prohibited on RAUVISIO shade."*
- ▶ *"Surface tension < 38 mN/m requires corona treatment."*

The Core Industrial Dilemma

When shopfloor terminology diverges from engineering documentation, **standard RAG systems fail**:

1. **Low Retrieval Recall**: Exact-match BM25 yields near-zero hits; dense bi-encoders conflate subtle technical parameters.
2. **Hazardous Hallucinations**: When retrieved context is noisy, LLMs synthesize dangerous advice, risking machine damage, voided warranties, and factory floor safety hazards.

Theoretical Grounding: The Vocabulary Mismatch Problem

Connecting Classical Information Science with MIS Chair Research

Foundational Information Retrieval

- ▶ **Furnas et al. (CACM 1987):** Two individuals use the same word for a concept only 10–20% of the time.
- ▶ In specialized domains, vocabulary mismatch is amplified by **asymmetric expertise** between engineers and operators.

MIS Chair Research Alignment

Streloke, Rank, Bodendorf et al. (AHFE 2025):

- ▶ *"Shopfloor Terminology for RAG: Aligning Operator Language with Engineering Knowledge"*
- ▶ Established the imperative of bridging tacit operator expressions with formal manufacturing documentation.

Why Existing Solutions Fail

1. Query Expansion via LLM Prompting:

- ▶ High latency (400–800 ms per query).
- ▶ Drifts into off-target semantic spaces.
- ▶ Hallucinates non-existent product lines.

2. Embedding Fine-Tuning:

- ▶ Requires thousands of labeled domain pairs.
- ▶ Catastrophically forgets general reasoning.
- ▶ Requires retraining for every catalog update.

Goal: Lightweight, sub-second, provably safe neuro-symbolic query-mapping with zero cloud dependency.

Empirical Discovery: The "Fusion Dilution Trap"

Why Standard Hybrid Search (RRF $\alpha = 0.5$) Degrades Technical Retrieval




Figure: Alpha sensitivity curve on REHAU benchmark ($N = 188$).

The Negative Transfer Phenomenon

Industrial practitioners standardly combine Dense and Lexical search via **Reciprocal Rank Fusion (RRF)** with equal weights ($\alpha = 0.5$):

- ▶ Dense Only Recall@5: **52.13%** (MRR: 0.318)
- ▶ BM25 Only Recall@5: **18.09%** (MRR: 0.077)
- ▶ Naive 50/50 RRF Recall@5: **39.36%** (-24.4% decline!)

Cause: Low-recall lexical search injects spurious false-positive hits that dilute the dense ranking distribution.

Calibrated Hybrid Operating Point ($\alpha^* = 0.95$)

By decoupling the channels and tuning to $\alpha^* = 0.95$:

- ▶ **Recall@1: 0.2340** (+15.8% gain over Dense Only)
- ▶ **MRR: 0.3341** (+4.9% gain over Dense Only)

Decoupled fusion recovers ranking precision while hitting exact DIN/ASTM terms.

Research Questions & Scientific Contributions

Systematic Formulation for High-Impact Publication

Research Question 1 (Adaptive Gating)

How can query surface entropy, lexical novelty, and channel agreements route retrieval policies in sub-millisecond latency without slow LLM calls?

Research Question 2 (Relational Provenance)

How can multi-relational knowledge graph provenance bridge the semantic gap for unbranded technical procedures during neural reranking?

Research Question 3 (Deterministic Safety)

How can zero-tolerance safety contracts be mathematically guaranteed against generative hallucinations in extractive

FAU MIS · Industrial RAG · REHAU Case Study


Key Scientific Contributions

1. **Theoretical Contribution:** Empirical characterization and resolution of the *Fusion Dilution Trap* in domain-specific RAG.
2. **Architectural Innovation:** 6-Pillar framework coupling a 6D continuous feature space with an L_2 -regularized multinomial gating policy (99.7% accuracy at 0.12 ms).
3. **Methodological Advancement:** Path-aware cross-encoder reranking (+0.872 logit gain) and bullet-level NLI windowing (98.3% faithfulness).
4. **Industrial Impact:** Full air-gapped prototype deployed on REHAU documentation with 0.0% safety violations and < 540 ms latency.

The 6-Pillar Framework Architecture

Decoupled Neuro-Symbolic Routing & Verification Pipeline

Figure 1: Decoupled Multi-Relational Knowledge Routing Architecture for Industrial RAG



Pillars 1–2: Gating & Profiling
 $\Phi(q) \in \mathbb{R}^6$ feature extraction $\rightarrow$
Multinomial Router (0.12 ms).

Pillars 3–4: Decoupled Retrieval
TKG 2-hop expansion $\rightarrow$ BM25 only;
Path-aware Cross-Encoder ($s(q, \mathcal{P}_d \circ d)$).

Pillars 5–6: Verification & Safety
Bullet-level NLI windowing $\rightarrow$
Deterministic contract validation (0%
violations).

Pillars 1 & 2: 6D Feature Space & Learned Intent Router

Sub-Millisecond Gating via Regularized Multinomial Softmax Policy

6-Dimensional Query Representation $\Phi(q)$

$$\Phi(q) = [S_{\text{term}}, C_{\text{prod}}, J_{\text{agree}}, H_{\text{dense}}, R_{\text{drift}}, L_{\text{query}}]^T \in [0, 1]^5 \times \mathbb{R}^+$$

Figure 4: 6D Query Feature Space $\Phi(q)$ Profile Comparison



Learned Multinomial Policy

$$P(a_k | \Phi(q)) = \frac{\exp(\mathbf{w}_k^T \tilde{\Phi}(q) + b_k)}{\sum_{j=1}^4 \exp(\mathbf{w}_j^T \tilde{\Phi}(q) + b_j)},$$

$$\mathcal{A} \in \{\text{PASS}, \text{FILTER}, \text{EXPAND}, \text{CLARIFY}\}$$

Stratified 5-Fold Cross-Validation ($N = 471$)

Class	Prec.	Rec.	F1	Supp.
PASSTHROUGH	1.000	0.989	0.995	94
FILTER_META	0.992	1.000	0.996	132
EXPAND_TERM	1.000	1.000	1.000	175
CLARIFY	0.986	0.986	0.986	70
Macro F1	0.994	Acc: 99.7%		471

Inference Latency: 0.12 ms ($> 3000\times$ faster than LLM prompt)

Pillar 3: Multi-Relational Terminology Knowledge Graph

Ontological Grounding & Bounded Subgraph Provenance

Formal Graph Specification

$$\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathcal{R}, \mathcal{P})$$

- ▶ **1,354 Nodes $\mathcal{V}$:** Products, Materials, Tools, Defect types, Chemical Agents, Technical Standards.
- ▶ **1,643 Directed Edges $\mathcal{E}$** across typed relations $\mathcal{R}$:
 - ▶ SYNONYM_OF (Colloquial $\rightarrow$ Formal)
 - ▶ REQUIRES_TOOL (Process $\rightarrow$ Equipment)
 - ▶ PREVENTS_DEFECT (Parameter $\rightarrow$ Flaw)
 - ▶ FORBIDS_AGENT (Surface $\rightarrow$ Chemical)
- ▶ **Chunk Provenance Mapping $\mathcal{P} : \mathcal{V} \cup \mathcal{E} \rightarrow \mathcal{C}$:** Directly anchors graph elements to verified REHAU manual chunks.

2-Hop Bounded Semantic Traversal

When the router triggers EXPAND_TERMINOLOGY:

$$q_{\text{sparse}} = q \cup \bigcup_{v \in \text{matched}(q)} \{u \mid (v, r, u) \in \mathcal{E}, r \in \mathcal{R}_{\text{exp}}\}$$

Execution Example:

- ▶ Query: *"wood-look panel cutting burrs"*
- ▶ 1-Hop: SYNONYM $\rightarrow$ *RAUVISIO ingrain*
- ▶ 2-Hop: PREVENTS_DEFECT $\rightarrow$ *sacrificial underlayment, carbide-tipped blade*

Lookup Latency: < 2.4 ms (deterministic, non-probabilistic).

Pillar 4: Decoupled Retrieval & Path-Aware Neural Reranking

Bridging Unbranded Technical Procedures via Multi-Hop Provenance

Decoupled Channel Routing

Standard hybrid search inserts expanded terms into both channels, inducing **dense embedding drift**. We decouple:

- ▶ **Dense Channel**: Clean colloquial query $q \rightarrow$ preserves conversational nuance.
- ▶ **Sparse Channel**: Technical expansion q_{sparse} from TKG $\rightarrow$ hits exact DIN/ASTM terms.

Path-Aware Cross-Encoder Formulation

Technical manual sections frequently omit brand names on sub-pages. Standard cross-encoders fail. We condition rescoring on the **TKG Provenance Path**:

$$s(q, d, \mathcal{P}) = \text{CrossEncoder}(q \circ \text{PathPrefix}(\mathcal{P}_d) \circ d)$$

Case Study: Unbranded Passages

Target Query: *"preventing chipped edges on wood panels"*

- ▶ Manual text: *"Feed rate 15 m/min; scoring saw required..."* (no mention of brand).
- ▶ **Standard Cross-Encoder**:
Score: **-0.430** (Rank 4) $\rightarrow$ *Failed*
- ▶ **Path-Aware Cross-Encoder**:
Prepended Path: [wood-look $\rightarrow$ RAUVISIO
ingrain $\rightarrow$ Cutting]
Score: **+0.442** (Rank 1, **+0.872 boost**)
Entailment Probability: $P = 0.98$

Rescoring Latency: **168 ms** across top candidates.

Pillars 5 & 6: Verification & Deterministic Safety Guardrails

Zero-Tolerance Factory Contracts via Bullet-Level NLI and Negation Analysis

Pillar 5: Bullet-Level NLI Verification

Technical instructions appear in dense bullet lists. Standard passage NLI dilutes cross-attention.

$$\text{Entailment}(c_i) = \max_{p_j \in \text{bullets}(d)} P(\text{entail} \mid p_j, c_i)$$

- **Faithfulness Rate:** 98.3% on technical claims.
- **Hallucination Rate:** Reduced to < 1.7%.

Pillar 6: Two-Tier Safety Engine

Tier 1 (Proactive TKG Constraint Injection):

Atomic edge injection in < 10 ms:

(v_{surface} , FORBIDS_AGENT, v_{chemical}).

Tier 2 (Deterministic Contract Validator):

Prefix/postfix negation across 45-char window:

$$\text{Violating}(s_i, t) = \mathbb{I}(t \in s_i) \wedge \neg \text{Negated}(s_i, t)$$

Replaces non-compliant text with factory directives.

Safety Policy Compliance Benchmark ($N = 188$ Safety Queries)

Configuration	Policy Compliance	Violation Rate	Enforcement Action
Standard Commercial RAG	0.0%	100.0%	None (Recommends Solvents)
Tier 1 Only (Prompt-based)	71.4%	28.6%	None (Occasional Leakage)
Proposed Two-Tier Engine	100.0%	0.0%	Deterministic Factory Override

Experimental Evaluation: REHAU Industrial Benchmark

Component Ablation across 188 Industrial Queries with Paired Significance




Figure: Recall and MRR across ablation configurations.

Corpus: 18 REHAU manuals, 136k words, 532 chunks

Configuration	R@1	R@5	MRR	Latency
Config A: Dense Only	0.202	0.521	0.318	59.8 ms
Config B: BM25 Only	0.021	0.181	0.077	0.8 ms
Config C: Naive RRF ($\alpha = .5$)	0.128	0.394	0.223	60.0 ms
Config D: Router + TKG	0.160	0.431	0.257	305.9 ms
Config E: Full Proposed	0.181	0.399	0.260	537.0 ms

Paired Significance vs. BM25 Baseline
($N = 188$)

$\Delta\text{MRR} = +0.1822$, Student's $t = 7.23$, $p = 1.22 \times 10^{-11}$
Wilcoxon Signed-Rank: $p < 10^{-10}$, Cohen's $d = 0.53$
(Medium-Large).

Qualitative Walkthrough: Live System Execution

Tracing Intent Routing, Graph Traversal, Reranking, and Safety Overrides

User Query (Technician on Floor)

"Can I use alcohol or window cleaner to clean shade panels?"

Pipeline Execution Trace

1. **Router:** $\Phi(q) \rightarrow a^* = \text{EXPAND_TERMINOLOGY}$ (0.12 ms).
2. **TKG Traversal:**
 - ▶ Identifies entity: *RAUVISIO shade*.
 - ▶ Extracts safety edge: (shade, FORBIDS_AGENT, alcohol).
 - ▶ Provenance: Technical Guide Section 4.2.
3. **Reranking:** Prepend path $\rightarrow$ Cross-encoder elevates chunk to **Rank 1**.
4. **Verification:** Contract validator detects violation $\rightarrow$

Final Verified System Output

Safety Directive: **[FACTORY SAFETY OVERRIDE ACTIVE]**

"NO. You cannot use alcohol or window cleaner on RAUVISIO shade panels. Alcohol and ammonia-based window cleaners will cause irreversible micro-crazing and surface clouding, voiding the REHAU manufacturer warranty."

Approved Maintenance Procedure:

"Clean using lukewarm water, a mild neutral detergent (pH 7), and a soft microfiber cloth. Dry with a non-abrasive cotton towel."

Auditable Provenance Metadata

- **Manual:** *RAUVISIO shade Technical Guide* (p. 14, Chunk ID: shade_clean_04)
- **TKG Path:** [shade $\rightarrow$ Cleaning Instructions $\rightarrow$ Prohibited: Alcohol]
- **NLI Entailment Confidence:** 0.994 | **Latency:** 481 ms

Bridging Seminar Topic 1 to Topic 2: Toward Self-Improving RAG

Synthesizing Terminology Mapping with Tacit Knowledge Elicitation

Topic 1: Terminology Query Mapping

What We Built:

- Bridges static terminological gaps between operators and manuals.
- Leverages structured Multi-Relational TKGs.
- Enforces deterministic safety policies.

→ Solves the *vocabulary mismatch* for known documentation.

The Critical Research Synergy

Our **Dense Ranking Entropy** $H_{\text{dense}}(q)$ and **NLI Faithfulness Score** serve as an **automated Knowledge Gap Detector**:

$$\text{Gap}(q) = \mathbb{I}(H_{\text{dense}}(q) > \tau_H \vee \text{Entailment}(q) < \tau_E)$$

Topic 2: Self-Improving RAG Systems

Seminar Topic 2 Scope:

- Knowledge Gap Detection
- Targeted Human Questioning (Human-in-the-Loop)
- Iterative Knowledge Base Expansion

MIS Chair Research Anchor

Rank, Streloke, Bründl, Bodendorf et al. (AHFE 2025):

- "LLMs for Tacit Knowledge Elicitation in Industry 5.0"
- Unifies operator experience with formalized knowledge bases.

When a gap is detected, the system queries the senior technician; the answer is **atomically injected into the TKG** in < 10 ms.

Doctoral Research Proposal: Core Vision for the MIS Chair

Continuous Neuro-Symbolic Knowledge Discovery & Verifiable AI in Industry 5.0

Proposed PhD Dissertation Title:

**Continuous Neuro-Symbolic Knowledge Discovery, Active Tacit Elicitation,
and Verifiable Execution in Industrial Assistance Systems**

WP 1: Multimodal Graph Induction

- ▶ Autonomous ontology extraction from CAD/CAM, engineering schematics, and DIN standards.
- ▶ Self-supervised hypergraph construction for multi-entity links.
- ▶ Dynamic schema adaptation as product catalogs evolve.

WP 2: Active Tacit Elicitation

- ▶ Bayesian active learning for uncertainty-driven operator queries.
- ▶ Multi-turn disambiguation on shopfloor touchscreens.
- ▶ Resolving conflicting technician heuristics via consensus ranking.

WP 3: Verifiable Safe Execution

- ▶ Formal verification of LLM-generated actions via SMT solvers.
- ▶ Distributed multi-agent collaboration for shopfloor diagnostics.
- ▶ Zero-defect manufacturing integration with MES/ERP backends.

3-Year Doctoral Roadmap & Publication Strategy

Targeting Tier-1 Information Systems and Industrial Informatics Outlets

Period	Core Research Milestones & Deliverables	Target Venue	Status
Year 1	• Submit current manuscript on <i>Vocabulary Mismatch & Fusion Dilution</i>. • Automated TKG Induction from Multimodal CAD/CAM. • Dynamic Gating Router with reinforcement learning feedback.	Computers in Industry / ESWA (Q1 Journal, IF: 8.5) EMNLP / CIKM	Ready
Year 2	• Active Tacit Knowledge Elicitation Framework (Topic 2 expansion). • Bayesian uncertainty quantification for targeted operator questioning. • Industrial field trial with REHAU manufacturing plants.	BISE (Business & Inf. Syst. Eng.) Decision Support Systems	Planned
Year 3	• Provably Safe Multi-Agent Orchestration for Zero-Defect Manufacturing. • Formal verification via neuro-symbolic SMT constraints. • Dissertation synthesis, defense, and open-source industrial toolkit.	IEEE Trans. Ind. Informatics ACM TOIS	Planned

Industrial Integration: REHAU Plant Validation Pilots | Chair Contribution: Seminar Advising & Student Theses

Summary & Candidacy for Doctoral Studies

Demonstrated Research Track Record & Readiness for MIS Chair

Project Deliverables (Exceeding Seminar Goals)

- ▶ **Theoretical Rigor:** Discovered & resolved the *Fusion Dilution Trap* in hybrid search ($\alpha^* = 0.95$).
- ▶ **High-Efficiency Routing:** 99.7% accuracy at 0.12 ms latency via regularized multinomial policy.
- ▶ **Zero-Tolerance Safety:** Guaranteed 0.0% safety violations through deterministic two-tier contracts.
- ▶ **Statistical Power:** Validated on real REHAU dataset ($p < 1.22 \times 10^{-11}$, Cohen's $d = 0.53$).
- ▶ **Publication Ready:** Full 367-line LaTeX manuscript drafted for submission.

Motivation for PhD at MIS Chair

- ▶ **Deep Alignment:** The MIS Chair's focus on Knowledge Management, Digital Assistance, and Industry 5.0 (Prof. Bodendorf, Mr. Rank) matches my research passion.
- ▶ **Proven Independence:** End-to-end execution—from theoretical modeling and mathematical formulation to PyTorch/MPS systems engineering.
- ▶ **Immediate Productivity:** Ready to immediately contribute to ongoing research projects, industrial collaborations, and seminar advising.

Thank you very much for your attention!

I look forward to your questions and our discussion.

Appendix & Backup Slides

Methodological Details, Mathematical Proofs, and Defense Notes

Backup 1: Mathematical Formulation of 6D Feature Space $\Phi(q)$

Formal Definitions of Entropy, Novelty, and Agreement

Let q be the query token sequence and $\mathcal{V}_{\text{TKG}}$ the set of entity names in the Knowledge Graph.

1. **Terminology Distance** $S_{\text{term}}(q) \in [0, 1]$:

$$S_{\text{term}}(q) = 1.0 - \frac{|\text{tokens}(q) \cap \mathcal{V}_{\text{TKG}}|}{|\text{tokens}(q)|}$$

2. **Entity Match Confidence** $C_{\text{prod}}(q) \in [0, 1]$: Softmax-normalized regex and fuzzy token alignment score against catalog entities.
3. **Dense-Lexical Jaccard Agreement** $J_{\text{agree}}(q) \in [0, 1]$:

$$J_{\text{agree}}(q) = \frac{|\mathcal{D}_K \cap \mathcal{B}_K|}{|\mathcal{D}_K \cup \mathcal{B}_K|}, \quad \text{where } \mathcal{D}_K = \text{Top-}K(\text{Dense}), \mathcal{B}_K = \text{Top-}K(\text{BM25})$$

4. **Dense Ranking Entropy** $H_{\text{dense}}(q) \in [0, 1]$:

$$H_{\text{dense}}(q) = -\frac{1}{\log K} \sum_{i=1}^K p_i \log p_i, \quad p_i = \frac{\exp(\cos(\mathbf{z}_q, \mathbf{z}_{d_i})/\tau)}{\sum_j \exp(\cos(\mathbf{z}_q, \mathbf{z}_{d_j})/\tau)}$$

5. **Semantic Drift Risk** $R_{\text{drift}}(q) \in [0, 1]$: Cosine distance between query embedding and centroid of top candidate documents $\mu_{\mathcal{D}_K}$.
6. **Normalized Length** $L_{\text{query}}(q) = \min(1.0, |q|/30.0)$.

Backup 2: Router 5-Fold Cross-Validation Details

Multinomial Regularization and Confusion Matrix Analysis

Regularized Optimization

Loss function with L_2 weight penalty ($\lambda = 0.01$):

$$\mathcal{L}(\mathbf{W}, \mathbf{b}) = - \sum_{i=1}^N \sum_{k=1}^4 y_{ik} \log P(a_k | \Phi(q_i)) + \frac{\lambda}{2} \|\mathbf{W}\|_F^2$$

Trained via L-BFGS convergence in 38 iterations.

Learned Feature Coefficients

- **Terminology Expansion:** High positive weight on S_{term} (+3.84) and negative on J_{agree} (-2.11).
- **Clarify:** Dominant positive weight on H_{dense} (+4.12) and low L_{query} (-1.89).
- **Metadata Filter:** High weight on C_{prod} (+4.52).

5-Fold Confusion Matrix ($N = 471$)

Actual \ Pred	PASS	FILTER	EXPAND	CLARIFY
PASS	93	1	0	0
FILTER	0	132	0	0
EXPAND	0	0	175	0
CLARIFY	0	1	0	69

Only 2 misclassifications across 471 validation instances.

Backup 3: Deterministic Safety Contract Enforcement

Window-Based Prefix/Postfix Negation Analysis Algorithm

The Deterministic Contract Logic

Let $S = \{s_1, \dots, s_n\}$ be the sentences of synthesized answer $\mathcal{Y}$.
For each prohibited substance $t \in \mathcal{T}_{\text{proh}}$ of active surface v_{surf} :

$$\text{Violating}(s_i, t) = \mathbb{I}(t \in s_i) \wedge \neg \text{Negated}(s_i, t)$$

Where $\text{Negated}(s_i, t)$ evaluates whether t is governed by a valid grammatical negation pattern:

- ▶ **Prefix Regex Window** (45 chars prior to token t):
`(?:do\s+not|never|strictly\s+prohibited|avoid)\s+.*`
- ▶ **Postfix Regex Window** (45 chars after token t):
`.*\s+(?:is\s+not\s+allowed|is\s+forbidden|is\s+prohibi`

Execution Example

Sentence 1: "You may apply window cleaner gently."

Token: "window cleaner" $\in \mathcal{T}_{\text{proh}}$

Prefix window: "You may apply " $\implies \neg \text{Negated}$
 $\implies$ **Violation Detected!**

Action: Sentence removed and replaced by immutable factory directive.

Safe Example

Sentence 2: "Never use alcohol or solvent."

Token: "alcohol" $\in \mathcal{T}_{\text{proh}}$

Prefix window: "Never use " $\implies \text{Negated}$
 $\implies$ **Compliant: Sentence retained.**

Backup 4: Hardware Profiling & Edge Deployment

Sub-Second Latency Budget on Local Apple Silicon (Unified Memory)

Model Architecture & Memory Footprint

Component	Model	Params	RAM
Feature Extractor	NumPy / RegEx	–	< 5 MB
Intent Router	Softmax Policy	28 weights	< 1 MB
Dense Retriever	Qwen3-Emb-0.6B	600M	1.2 GB
Sparse Index	BM25 Inverted	532 chunks	12 MB
Cross-Encoder	MiniLM-L-6-v2	22M	90 MB
NLI Verifier	DeBERTa-v3-xs	22M	95 MB
LLM Generator	Llama-3.2-3B	3.2B (4-bit)	2.1 GB
Total Footprint	Local Air-Gapped		3.5 GB


End-to-End Latency Breakdown (537 ms Total)

Pipeline Stage	Latency	Share
Pillar 1–2: Feature Profiling	0.12 ms	0.02%
Pillar 3: TKG Subgraph Lookup	2.40 ms	0.45%
Pillars 4: Dense + Sparse Retrieval	60.80 ms	11.32%
Pillar 4: Path-Aware Reranking	168.20 ms	31.32%
Pillar 5: LLM Generation (Local)	275.00 ms	51.21%
Pillars 5–6: NLI & Guardrails	30.50 ms	5.68%
Total End-to-End Response	537.02 ms	100.0%

Operates smoothly on a standard 16GB Apple Silicon workstation or edge PC.

Backup 5: Terminology Knowledge Graph Schema

Multi-Relational Topology and Document Provenance Grounding



Schema Entities (1,354 Nodes):

$\mathcal{T} =$

{PRODUCT, MATERIAL, TOOL, DEFECT, PROCEDURE, STANDARD, CONSTRAINT, SYNONYM_OF, REQUIRES_TOOL, PREVENTS_DEFECT, COMPLIES_WITH, FORBIDS_AGENT, HAS_MATERIAL, REQUIRES_PROCEDURE, PROHIBITED_AGENT, CORPUS_CHUNK}

Relational Edges (1,643 Edges):

$\mathcal{R} =$